



Object Recognition

Objectives

To learn and appreciate the following concepts:

- ✓ Introduction to Support Vector Machines
- ✓ Face detection & Recognition
- ✓ Bag of Words
- ✓ Deep Learning



Session outcome

- At the end of session the student will be able to understand:
 - What is SVM
 - Fundamentals of Face detection and recognition
 - Importance of Bag of words
 - Understanding if deep learning

Support Vector Machine (SVM)

- Support Vector Machines (SVMs) are powerful supervised learning models used for classification and regression tasks. The primary objective of SVM is to find the optimal hyperplane that best separates data points into different classes.
- Support Vector Machines are powerful classifiers that work well in high-dimensional spaces and provide robustness against outliers. However, they can be computationally intensive and may have difficulty with interpretability, especially in complex scenarios.

How SVM Works:

- 1. Hyperplane:** In a binary classification scenario, a hyperplane is a line that separates data points into two classes. In higher dimensions, this hyperplane becomes a plane or a hyperplane.
- 2. Support Vectors:** These are the data points closest to the hyperplane. They are crucial because they define the decision boundary and the margin. SVM aims to maximize the margin, which is the distance between the hyperplane and the nearest data points.
- 3. Kernel Trick:** SVMs can efficiently perform classification even in high-dimensional spaces by implicitly mapping inputs into higher-dimensional feature spaces. The kernel trick allows SVMs to compute the dot product of vectors in these higher-dimensional spaces without explicitly calculating the transformation.



Support Vector Machine (SVM)

- When dealing with two classes that are not linearly separable, we can still use a linear Support Vector Machine (SVM) by employing techniques such as kernel trick or soft margin SVM. Here's a detailed approach to using a linear SVM for such scenarios:

1. Understanding the Data:

- First, it's essential to understand the nature of the data and why it's not linearly separable. Visualizing the data can provide insights into its distribution and the relationship between features and classes.

2. Kernel Trick:

- If the data is not linearly separable, we can map it to a higher-dimensional space where it might become linearly separable. This is achieved through the kernel trick, which implicitly maps the data into a higher-dimensional space without explicitly computing the transformation.

3. Kernel Selection:

- For a linear SVM, we typically use a linear kernel. However, if the data is not linearly separable, we can try other kernel functions such as polynomial, radial basis function (RBF), or sigmoid to map the data into a higher-dimensional space.

4. Soft Margin SVM:

- In cases where the data is not perfectly separable, we can relax the constraint and allow for some misclassification. This is done using a soft margin SVM, which introduces a penalty for misclassification in the objective function. The parameter C controls the trade-off between maximizing the margin and minimizing the classification error.

5. Cross-Validation:

- To determine the optimal parameters (such as the choice of kernel and the value of C), we can use techniques like cross-validation. This involves splitting the data into training and validation sets multiple times and evaluating the model's performance with different parameter values.

Support Vector Machine (SVM) - Example

- Let's consider a simple example of classifying whether a fruit is an apple or an orange based on two features: size and color.
- **Features:** Size (X-axis), Color (Y-axis)
- **Data Points:** Each data point represents a fruit, where apples are marked with blue circles and oranges with red triangles.
- In this example, the solid line represents the decision boundary (hyperplane) determined by the SVM. The dashed lines on either side represent the margins. The support vectors, marked by larger markers, are the data points closest to the decision boundary.
- Using a linear SVM for non-linearly separable data involves techniques like kernel trick and soft margin SVM. By mapping the data into a higher-dimensional space and allowing for some misclassification, we can effectively classify non-linearly separable classes.

Support Vector Machine (SVM) – Advantages & Disadvantages

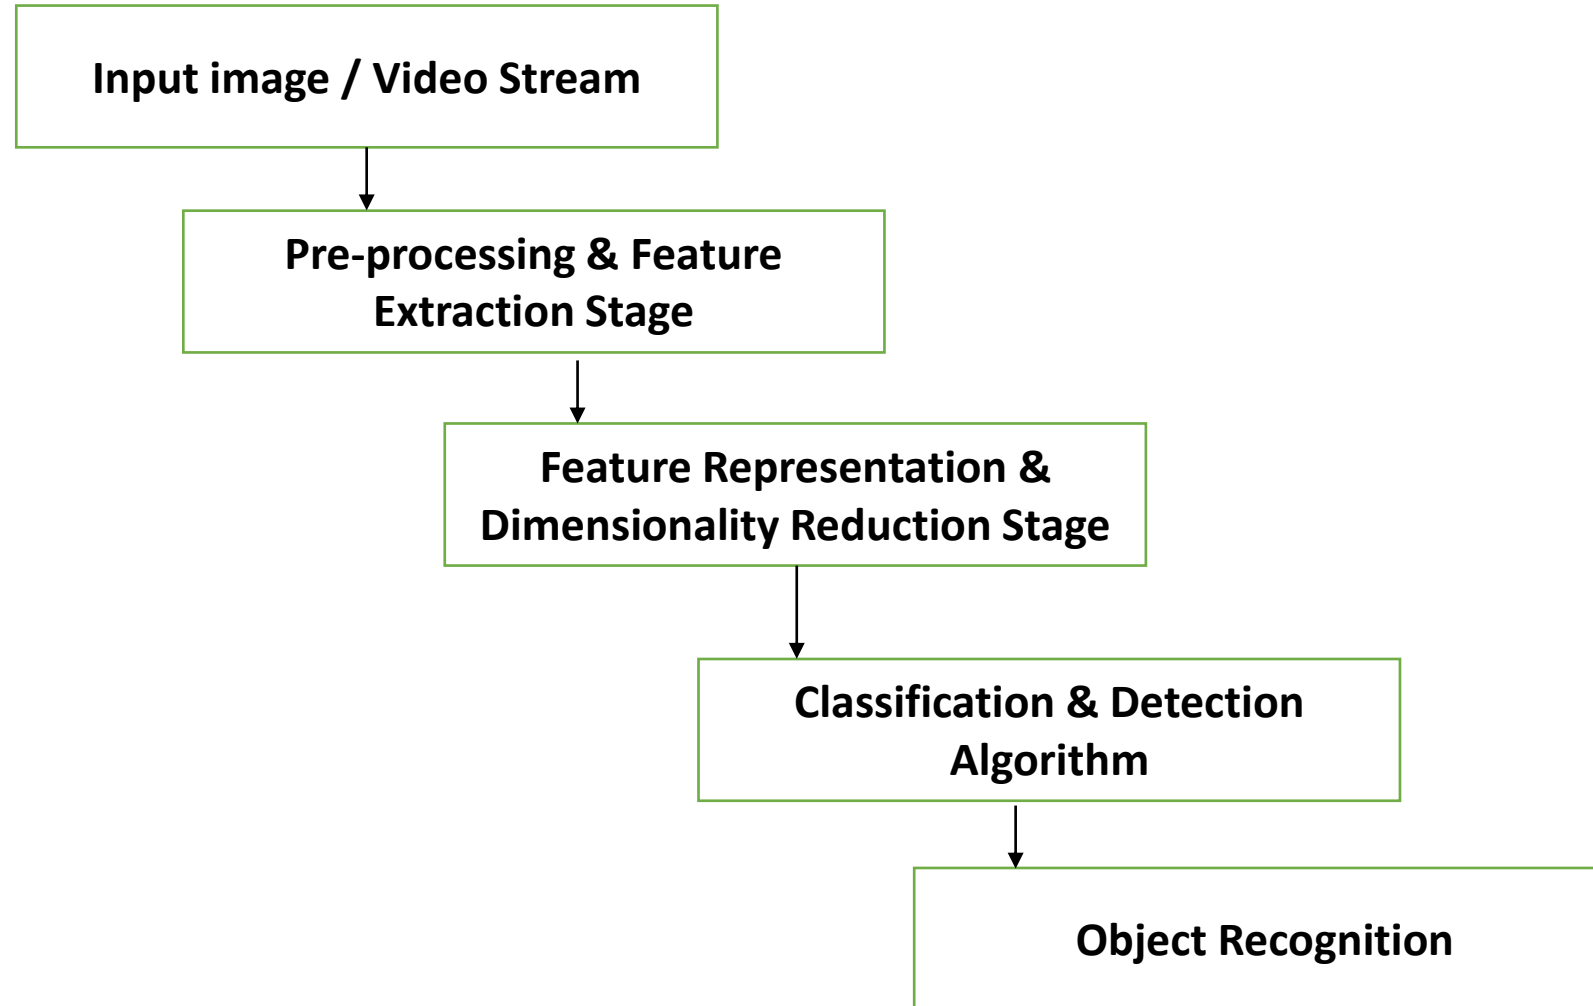
Advantages of SVM:

- 1. Effective in High-Dimensional Spaces:** SVMs perform well even when the number of dimensions exceeds the number of samples.
- 2. Versatile:** SVMs can be used for both classification and regression tasks.
- 3. Robustness:** SVMs are effective in handling outliers due to the use of support vectors.
- 4. Memory Efficient:** SVMs only use a subset of training points in the decision function, making them memory efficient.

Disadvantages of SVM:

- 1. Computationally Intensive:** Training an SVM can be computationally expensive, especially with large datasets.
- 2. Difficulty in Interpretability:** The hyperplane generated by SVM might be hard to interpret, especially in higher dimensions.
- 3. Sensitive to Noise:** SVMs can be sensitive to noisy data, leading to overfitting if not properly regularized.

Block Diagram Of A Simple Pipeline For Object Recognition





Block Diagram Of A Simple Pipeline For Object Recognition

In this block diagram:

Step 1: The pipeline starts with an input image or video stream.

Step 2: The Pre-processing & Feature Extraction Stage involves operations like image resizing, normalization, and extracting relevant features from the input data.

Step 3: The Feature Representation & Dimensionality Reduction Stage transforms the extracted features into a suitable representation and reduces their dimensionality to improve efficiency.

Step 4: The Classification or Detection Algorithm applies machine learning or computer vision techniques to classify objects or detect their presence in the input data.

Step 5: Finally, the Object Recognition stage identifies and labels the recognized objects.

Few Challenges For Existing Visual Recognition Algorithms

1. **Variability in Object Appearance:** Objects can appear differently due to changes in illumination, viewpoint, occlusion, scale, and deformation, making it challenging for algorithms to recognize them accurately under different conditions.
2. **Large-Scale Data Handling:** Processing large-scale datasets with millions of images requires efficient algorithms and computational resources.
3. **Limited Training Data:** Training deep learning models requires a large amount of labeled data, which may not always be available, especially for rare or fine-grained categories.
4. **Overfitting:** Deep learning models are prone to overfitting, where they memorize training examples instead of learning generalizable patterns, leading to poor performance on unseen data.
5. **Computational Complexity:** Deep learning models, especially convolutional neural networks (CNNs), are computationally intensive and require powerful hardware accelerators for real-time performance in applications like autonomous vehicles and robotics.
6. **Interpretability:** Deep learning models are often considered black boxes, making it difficult to understand the reasoning behind their predictions, which is crucial for building trust in AI systems.
7. **Domain Adaptation:** Models trained on one dataset may not generalize well to other datasets due to domain shifts, necessitating techniques for domain adaptation to improve performance across different domains.



Bag of Words

- Refer PDF

Deep Learning

- There are several deep learning architectures commonly used for object recognition tasks.
 - These architectures have been used in various object recognition tasks, including image classification, object detection, semantic segmentation, instance segmentation, and more. The choice of architecture depends on factors such as the nature of the task, available computational resources, and the trade-off between accuracy and efficiency.
1. **Convolutional Neural Networks (CNNs):** CNNs are one of the most widely used architectures for object recognition tasks. They consist of multiple layers of convolutional and pooling operations followed by fully connected layers. Examples include:
 - AlexNet
 - VGG (Visual Geometry Group) networks
 - GoogLeNet (Inception)
 - ResNet (Residual Network)
 - DenseNet
 - MobileNet
 2. **Region-Based Convolutional Neural Networks (R-CNN):** R-CNN and its variants are popular for object detection tasks where the goal is to localize objects within an image. Examples include:
 - Fast R-CNN
 - Faster R-CNN
 - Mask R-CNN

Deep Learning

- 3. Single Shot MultiBox Detector (SSD):** SSD is a real-time object detection framework that predicts object classes and bounding boxes directly from feature maps of different scales.
- 4. You Only Look Once (YOLO):** YOLO is another real-time object detection system that processes images in a single pass and predicts bounding boxes and class probabilities simultaneously.
- 5. Capsule Networks:** Capsule networks aim to overcome some limitations of CNNs by encoding both the presence of features and their spatial relationships. Examples include CapsuleNet.
- 6. Attention Mechanisms:** Attention mechanisms have been incorporated into deep learning architectures to selectively focus on relevant parts of an image or sequence of features. Examples include Transformer-based architectures like Vision Transformer (ViT).
- 7. Graph Neural Networks (GNNs):** GNNs are used for object recognition tasks where the input data can be represented as graphs, such as scene graphs or point clouds.
- 8. Siamese Networks:** Siamese networks are used for tasks like similarity learning and one-shot learning, where the goal is to determine if two input images are of the same class.